

# Fast, Easy Linear Algebra in C++ with Eigen

Benjamin S. Skrainka

University of Chicago  
The Harris School of Public Policy  
[skrainka@uchicago.edu](mailto:skrainka@uchicago.edu)

May 16, 2012

# Objectives

This talk's objectives:

- ▶ State benefits of using Eigen from linear algebra
- ▶ List key components of Eigen
- ▶ Use Eigen to perform typical linear algebra operations
- ▶ Generate high quality pseudo-random numbers with Mersenne Twister

# Plan for this Talk

Getting Started

Eigen's Containers:  $\text{Matrix}<T>$  and  $\text{Array}<T>$

Advanced Operations

Mersenne Twister

# Benefits of Eigen

Eigen makes dense linear algebra easy in C++:

- ▶ All basic features and many more advanced ones as well
- ▶ Template-based  $\Rightarrow$  works with all arithmetic types
- ▶ Syntax is almost as easy to use as MATLAB's
- ▶ Blazing fast performance because of expression templates:
  - ▶ 10-30x speed up over MATLAB
  - ▶ Comparable to Intel MKL, the gold standard
  - ▶ Produces efficient code so arrays are only traversed once!
- ▶ Easy to install
- ▶ Good documentation
- ▶ For a single core
- ▶ Use PETSc or equivalent for MPI-based linear algebra on multiple cores

# Eigen Benchmark

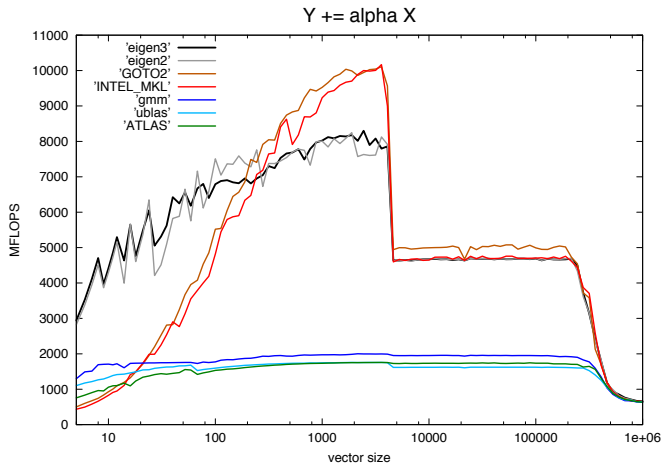


Figure is taken from Eigen documentation.

# Eigen Benchmark

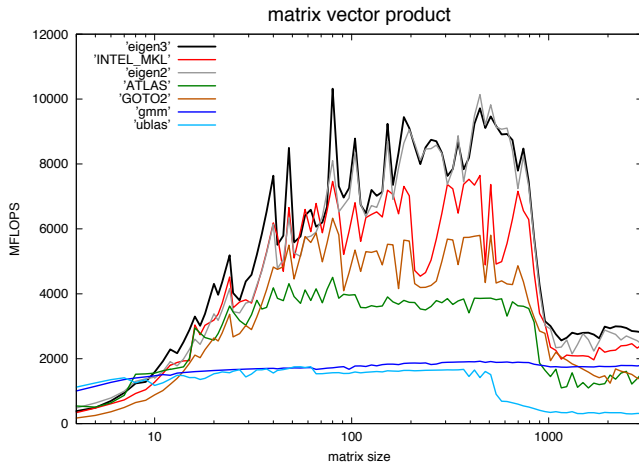


Figure is taken from Eigen documentation.

# Eigen Benchmark




Figure is taken from Eigen documentation.

# Overview of Eigen Libraries

Eigen provides several templated Matrix and Array types:

- ▶ `Matrix<T>`: generic matrix class
- ▶ `MatrixXd`: matrix of dynamically allocated doubles
- ▶ `VectorXd`: column vector of dynamically allocated doubles
- ▶ `RowVectorXd`: row vector of dynamically allocated doubles
- ▶ Plus common fixed sizes and types: `Matrix3f`, `Matrix2i`, etc.
- ▶ `Map<T>`: provides `Matrix<T>` interface to an existing slab of data

`Matrix<T>` operations are matrix arithmetic (like `*` in MATLAB) whereas `Array<T>` operations are coefficient-wise (like `.*` in MATLAB).



# Math Operations

Eigen supports all the standard math functions via member functions:

- Includes:

<code>abs()</code>	<code>pow()</code>
<code>min()</code>	<code>exp()</code>
<code>max()</code>	<code>log()</code>
<code>square()</code>	<code>sqrt()</code>

- Must convert `Matrix<T>` to `Array<T>` using `array()` member function in order to perform coefficient-wise operations
- Performed coefficient-wise, i.e. like MATLAB where

`m.array().exp()` is  $\begin{pmatrix} \exp(a_{12}) & \exp(a_{12}) \\ \exp(a_{21}) & \exp(a_{22}) \end{pmatrix}$

# Overview of Basic Linear Solving

Eigen also contains many linear solvers:

Decomposition	Speed	Accuracy
PartialPivLU	++	+
FullPivLU	-	+++
HouseholderQR	++	+
ColPivHouseholderQR	+	++
FullPivHouseholderQR	-	+++
LLT (Cholesky)	+++	+
LDLT (Robust Cholesky)	+++	++

There are also SVD and eigenvalue decompositions

# Array<T> vs. Matrix<T>

Select Array<T> or Matrix<T> depending on desired behavior:

- ▶ Matrix<T>:
  - ▶ Designed for linear algebra
  - ▶ Supports matrix arithmetic
  - ▶ Has some coefficient-wise member functions
  - ▶ Convert to an Array<T> with Matrix<T>::array()
- ▶ Array<T>:
  - ▶ General purpose array
  - ▶ Coefficient-wise arithmetic
  - ▶ Can perform coefficient-wise operations, such as add a number to all elements
  - ▶ Convert to Matrix<T> with Array<T>::matrix()

# Getting Started

# First Steps

Let's start by discussing:

- ▶ Installation
- ▶ Configuration
- ▶ Getting Help

# Installation

Eigen is incredibly easy to install:

- ▶ Entire library consists of header files
- ▶ No need to compile or build any binary libraries
- ▶ Just download and copy it to where you want it

To install,

- ▶ Download latest stable release from  
<http://bitbucket.org/eigen/eigen/get/3.0.5.tar.gz>
- ▶ Unpack:

```
$ tar xzvf eigen-eigen-6e7488e20373.tar.gz
```

```
$ mv eigen-eigen-6e7488e20373 eigen-3.0.5
```

# Help

There are many resources for help:

- ▶ Online [documentation](#)
- ▶ Online [Getting Started Tutorial](#)
- ▶ Online [Quick Reference](#)

# A First Application

```
#include <iostream>
#include <Eigen/Dense>

int main()
{
    Eigen::MatrixXd m(2,2) ;
    m(0,0) = 3 ;
    m(1,0) = 2.5 ;
    m(0,1) = -1 ;
    m(1,1) = m(1,0) + m(0,1) ;
    std::cout << m << std::endl ;

    m.diagonal() << -1, -2 ;
    std::cout << m < std::endl ;
    return 0 ;
}
```



# Building an Application with Eigen

To build an application with Eigen, use a supported C++ compiler such as GNU or MSVC++:

- ▶ Specify path to Eigen header files with `-I/path/to/Eigen`
- ▶ Do not compile with `-pedantic`:
  - ▶ `-pedantic` forces strict conformance to ANSI/ISO C++ standards
  - ▶ But, Eigen requires `long long` which is not yet officially supported...
- ▶ Disable debugging checks in product code:
  - ▶ Compile with `-DEIGEN_NO_DEBUG`
  - ▶ Huge speed up!

# Example Compilation

Here is an example compilation

- For debugging (development):

```
$ g++ -Wall -Wextra -O0 -g -m64 \
-I/home/skrainka/public/sw/include \
eigenExample.cpp -o eigenExample
```

- For production code:

```
$ g++ -Wall -Wextra -O2 -m64 \
-DEIGEN_NO_DEBUG \
-I/home/skrainka/public/sw/include \
eigenExample.cpp -o eigenExample
```

## Practical Tips on Use

For production code, compile with `-DEIGEN_NO_DEBUG`:

- ▶ Removes checks for conformability, etc.
- ▶ Results in much faster code (10-30x!)
- ▶ Make sure your code is debugged first!

Use typedefs:

- ▶ For floating point type
- ▶ For key matrix and vector types

BLP Example:

```
typedef long double BLPReal_t ;  
typedef Eigen::Matrix< BLPReal_t, Eigen::Dynamic,  
    Eigen::Dynamic, Eigen::ColMajor | Eigen::AutoAlign >  
    BLPMatrixXd ;  
typedef Eigen::Matrix< BLPReal_t, Eigen::Dynamic, 1,  
    Eigen::ColMajor | Eigen::AutoAlign > BLPVectorXd ;
```

# Configuration Options

Eigen is highly configurable using compile-time macros. Some options are:

- ▶ `EIGEN_DEFAULT_TO_ROW_MAJOR`
- ▶ `EIGEN_INITIALIZE_MATRICES_BY_ZERO`
- ▶ `EIGEN_NO_DEBUG`

See the manual for details.

# Eigen's Containers: $\text{Matrix}<T>$ and $\text{Array}<T>$

# Matrix<T>

`Matrix< typename Scalar, int nRows, int nCols >:`

- ▶ Scalar: data type to store in Matrix (double, int, etc.)
- ▶ nRows and nCols:
  - ▶ Number of rows and columns
  - ▶ Pre-specified number of rows and columns at compile time  $\Rightarrow$  'fixed size'
  - ▶ Can use `Eigen::Dynamic` if size is unknown  $\Rightarrow$  'dynamic' size
  - ▶ Use fixed size for smaller matrices – i.e. size  $\leq 32$
- ▶ Can also set `Eigen::RowMajor` or `Eigen::ColMajor`

# Vectors

A vector is just a Matrix:

- ▶ A column vector is

```
Eigen::Matrix< typename Scale, int nRows, 1 >
```

- ▶ A row vector is

```
Eigen::Matrix< typename Scale, 1, int nCols >
```

- ▶ Orientation matters:

- ▶ VectorXd is a column vector
- ▶ RowVectorXd is a row vector

# Predefined Types

Eigen provides several predefined types:

- ▶ **MatrixNt**
- ▶ **VectorNt**
- ▶ **RowVectorNt**
- ▶ where

<b>N</b> : Storage	<b>t</b> : type
<b>X</b> : dynamic	<b>d</b> : double
<b>n</b> : number of rows and/or columns (2, 3, or 4)	<b>f</b> : float
	<b>i</b> : int
	<b>cd</b> : complex double
	<b>cf</b> : complex float

- ▶ Examples:
  - ▶ **MatrixXd**: dynamically allocated matrix of doubles
  - ▶ **Vector3i**: vector with three ints



# Coefficient Accessors

Element access is intuitive:

- ▶ Access elements using operator():

```
Matrix3i mMat ;  
mMat( 1, 2 ) = 3 ;
```

- ▶ Remember: 0-based indexing!
- ▶ Can also access sub-matrices (slices) like the ':' operator in MATLAB

# Initialization

Initialization is easy and intuitive:

- ▶ Use operator<< to fill by column then row:

```
Matrix3d m3 ;  
m3 << 1, 2, 3,  
      4, 5, 6,  
      7, 8, 9 ;
```

- ▶ Can use member functions
  - ▶ setZero()
  - ▶ setOnes()
  - ▶ setConstant()
  - ▶ setRandom()
  - ▶ setLinSpaced()
- ▶ See manual for more complex options

# Matrix Arithmetic

Operators are overloaded so that arithmetic behaves as expected:

```
Eigen::Matrix3d m1, m2 ;  
Eigen::Vector3d v1, v2 ;  
  
v2 = m1 * v1 ; // Matrix multiplication  
m2 = 5.0 * m1 ; // Scalar multiplication  
v1 += v2 ;  
m1 = m2.transpose() ;
```

# Handy Member Functions

`Matrix<T>` has several handy member functions:

- ▶ Ctors:
  - ▶ `MatrixNt( nRows, nCols )`
  - ▶ `VectorNt( nElem )`
- ▶ `m.size()` - number of elements
- ▶ `m.cols()` - number of columns
- ▶ `m.rows()` - number of rows
- ▶ `m.resize()` - resize matrix

# Reductions

Eigen provides several convenient matrix reduction operations:

<code>m.minCoeff()</code>	minimum coefficient
<code>m.maxCoeff()</code>	maximum coefficient
<code>m.prod()</code>	product of coefficients
<code>m.sum()</code>	sum of coefficients
<code>m.mean()</code>	mean of coefficients
<code>m.trace()</code>	trace of matrix
<code>m.colwise()</code>	perform reduction on each column
<code>m.rowwise()</code>	perform reduction on each row
<code>m.all()</code>	true iff all elements are true
<code>m.any()</code>	true iff any element is true
<code>m.lpNorm&lt; p &gt;()</code>	compute the $\ell^p$ norm
<code>m.lpNorm&lt; Infinity &gt;()</code>	compute the $\ell^\infty$ norm

## Sub-Matrices & Sub-Vectors

Matrix<T>		
col()	topLeftCorner()	topRightCorner()
row()	bottomLeftCorner()	bottomRightCorner()
block()	topRows()	leftCols()
diagonal()	bottomRows()	rightCols()

Vector<T>
head()
tail()
segment()

## Array<T>

Array<T> has same template parameters as Matrix:

- ▶ Array<typename Scalar, int nRows, int nCols >
- ▶ Use for coefficient-wise operations
- ▶ Use Matrix<T> for linear algebra!
- ▶ \* and / are like .\* or ./ in MATLAB
- ▶ Use m.array() and a.matrix() to convert back and forth
- ▶ Sometimes you need both:

```
MatrixXf m(2,2);  
MatrixXf n(2,2);  
MatrixXf result(2,2);  
m << 1,2,  
      3,4 ;  
n << 5,6,  
      7,8 ;  
result = (m.array() + 4).matrix() * m ;  
result = (m.array() * n.array()).matrix() * m ;
```

# Coefficient-Wise Operations

Can often avoid converting a `Matrix<T>` to an `Array<T>` by using coefficient-wise operations which operate on two matrices:

- ▶ `m.cwiseMin( m2 )`
- ▶ `m.cwiseMax( m2 )`
- ▶ `m.cwiseProduct( m2 )`
- ▶ `m.cwiseQuotient( m2 )`
- ▶ `m.cwiseAbs( m2 )`
- ▶ `m.cwiseInverse( m2 )`
- ▶ `m.cwiseSqrt( m2 )`

See manual for details



# Advanced Operations

# Overview of Common Operations

Overview of common operations:

- ▶ Solving  $Ax = b$ :
  - ▶ Do not compute the inverse of a matrix unless you need it!
  - ▶ Solve in place if possible
  - ▶ Choose an algorithm with the appropriate speed and accuracy
  - ▶ Can use `m.inverse()` to compute inverse for small matrices
- ▶ Common decompositions and algorithms:
  - ▶ LU
  - ▶ QR
  - ▶ Cholesky
  - ▶ Eigenvalues
- ▶ Check for errors or numerical problems such as ill-conditioning

# Basic Linear Solving

To solve  $Ax = b$  for  $x$ :

1. Choose a suitable solver:

```
Matrix3d A ;  
Vector3d b ;  
...  
FullPivLU< Matrix3d > lu( A ) ;
```

Can also access most decompositions via a member function.

2. Solve the problem:

```
Vector3d x = lu.solve( b ) ;
```

3. Check that solver was successful:

```
if( lu.rank() != A.cols() )  
    exit( -1 ) ;
```

# Checking Linear Solution

The different decompositions provide several member functions which you can use to check if a solution was successful:

- ▶ `m.rank()` - use to check solution is full rank
- ▶ `m.setThreshold()` - sets threshold on rank-revealing algorithms to determine rank

Can also compute the relative error using the  $\ell^2$  `norm()` member function:

```
double dwRelErr = ( A * x - b ).norm() / b.norm() ;
```

# Least Squares

To solve least squares problems:

- ▶ SVD is usually best
- ▶ LDLT can be faster but less reliable

```
MatrixXf A = MatrixXf::Random(3, 2) ;  
cout << "A:\n" << A << endl ;  
VectorXf b = VectorXf::Random(3) ;  
cout << "b:\n" << b << endl ;  
cout << "Solution:\n"  
    << A.jacobiSvd( ComputeThinU | ComputeThinV ).solve(b)  
    << endl ;
```

# Include Files

Can optimize build time by choosing just the right header files:

- ▶ Easiest to `#include "Eigen/Dense"` for dense linear algebra
- ▶ May wish to choose header files with finer granularity to accelerate build times and use less memory:
  - ▶ `#include "Eigen/Core"`
  - ▶ `#include` other Eigen header files as appropriate
- ▶ See manual for details

# Map<T>

Sometimes you have data from a source external to Eigen (such as a solver) and it does not make sense to copy data into a `Matrix<T>`:

- ▶ External library expects an array, such as a solver
- ▶ But also need to perform linear algebra
- ▶ Use `Eigen::Map<T>` to create an Eigen wrapper for accessing the data as if it were an `Eigen::Matrix<T>`:

```
double vJac[ 16 ] ;
```

```
...
```

```
Eigen::Map< Eigen::Matrix4d > mJac( vJac ) ;
```

```
Eigen::Map< Eigen::MatrixXd > mJac( vJac, 16 ) ;
```

- ▶ Make sure you have correct ordering – column major or row major!

## Map<T> Example

```
integer n  = nTheta1 + nTheta2 + nShares + nIV ;  
double *x  = double[n] ;  
Eigen::Map< Eigen::Matrix< double,  
              Eigen::Dynamic, 1 > > vXOptTmp( x, n ) ;
```



## Map<T> Example II

```
int nElem = 3 ;
int m3i[ nElem ][ nElem ] ;
int count = 1 ;
for( int ix = 0 ; ix < nElem * nElem ; ++ix )
{
    for( int jx = 0 ; jx < nElem ; ++jx )
    {
        m3i[ ix ][ jx ] = count++ ;
    }
}
Eigen::Map<
    Eigen::Matrix<int, Eigen::Dynamic, Eigen::Dynamic,
                  Eigen::RowMajor >
    > mMap( &(m3i[0][0]), nElem , nElem ) ;
std::cout << "mMap:\n" << mMap << std::endl ;
```

# Other Features

Eigen has several other features you may need:

- ▶ `cast<T>` member function casts all elements in a matrix:

```
BLPVectorXd vXOpt = vXOptTmp.cast< blp::BLPReal_t >() ;
```

- ▶ `SparseMatrix`:
  - ▶ Handy for representing matrices which are large and mostly zero
  - ▶ Eigen provides some sparse solvers as well
- ▶ See the manual for details

# Performance Enhancements

To maximize performance:

- ▶ Compile without debug symbols and enable -O2
- ▶ Disable conformance checking via -DEIGEN\_NO\_DEBUG
- ▶ Use `.noalias()` if RHS can be safely evaluated into LHS without aliasing:

```
c.noalias() += a * b ;
```

- ▶ Write matrix expressions to favor either row major or column major ordering
- ▶ Initialize matrix and use member functions with fixed sizes specified at compile time for `m.size() <= 32`.

# Mersenne Twister

# Random Numbers

Generating good random numbers is necessary for research:

- ▶ Monte Carlo studies
- ▶ Synthetic Data
- ▶ Starting guesses

Machine-generated random numbers are *pseudo*-random:

- ▶ Not truly random
- ▶ Pass some set of statistical tests
- ▶ Mersenne Twister is currently the best algorithm
- ▶ M. Matsumoto and T. Nishimura, *ACM Transactions on Modeling and Computer Simulation*, Vol. 8, No. 1, January 1998.

# Installation

Mersenne Twister is easy to install:

1. Download `Mersenne-1.1.tar.gz` from the course website
2. Unpack with tar:

```
tar xzf Mersenne-1.1.tar.gz
```

3. `#include "MersenneTwister.h"` in your code – there is no need to build any libraries

This implementation is by Richard J. Wagner

# Documentation

There is no documentation:

- ▶ See README file in installation
- ▶ See `example.cpp` file
- ▶ See `MersenneTwister.h` header file

# Use

This implementation is very easy to use:

1. Instantiate an MTRand object
2. Use member functions to generate random numbers on desired interval (open, closed; normal, uniform)

Recommended usage:

```
#include "MersenneTwister.h"
...
MTRand mtrand;    // automatically generate seed
double s = mtrand();           // in [0,1]
double t = mtrand.randExc(0.5); // in [0,0.5)
unsigned long u = mtrand.randInt(10); // in [0,10]
```



## Other Issues

You should always set a seed so your work is reproducible:

```
MTRand mtrand( 1945 ) ; // Seed with 1945
```

- ▶ Test that results do not depend on the seed you have chosen.
- ▶ Do not use the system clock as a seed. It is not random...
- ▶ See `example.cpp` for using a more complex seed

Can generate Normally distributed random numbers:

```
mtrand.randNorm(  $\mu$ ,  $\sigma$  ) ;
```

Other features:

- ▶ Can save and load streams of random numbers
- ▶ Can scale random numbers for desired interval